

Ghost Predator — Patch Log, June 11–12 2026

Every change deployed in the last 36 hours. All items below are LIVE and verified working as of June 12, 13:45 UTC. Each shipped with: timestamped backup, syntax check, restart between windows, live verification.

1 · Entry timing — the sweet-spot window (the big one)

Deployed Jun 11, 13:05 UTC. Fire window changed from 'first qualifying tick of the final 30s' (fires landed ~T-29s) to **only the final 7–14 seconds** (SNIPE_WINDOW_SECS 30→14, MIN_FILL_SECS 2→7, both durations). Three independent datasets agreed: (a) 2,866-window tape study — a 1.5–3bps leader holds 74.8% at T-30 but 85.5% at T-10, while the in-band ask does NOT rise while waiting (the 0.75 cap filters priced-out favorites); (b) all 790 historical trades bucketed by seconds-left at fire — everything fired with 15+ seconds left totaled –\$220, fires under 15s were +\$41 (the only profitable bucket); (c) a 1,688-window replay with real asks — T-29..21 entries ran –\$0.24 to –\$1.00 EV per trade, T-15..7 ran positive (up to +\$1.59). Same shape on 527 15m windows.

Results since deploy: 19 trades, 15W–4L (79%), +\$35.54, avg EV +\$1.87/trade vs –\$0.39 in the old era. The bot trades roughly half as often and refuses windows the crowd has already priced (asks 0.91–0.99).

2 · Risk & guardrails

\$240 hard balance floor (Jun 11, 13:05): halts all trading if realized balance minus in-flight stakes would cross \$240 — checked before every order. Calibration drift fixed the same hour: the bot's internal balance had drifted ~\$30 above real exchange cash (accumulated fees); STARTING_BALANCE recalibrated so the floor counts real dollars.

Rolling-24h limit removed (Jun 11, 15:03): it halted the bot twice with zero new losses — yesterday's wins aging out of the trailing window deepened the rolling sum on time passing alone. Halt stack is now: \$240 floor + \$150 daily loss limit + 10-loss streak + max-open-positions.

Weekend entry bar: 5.0bps (set Fri 01:00 ET) → 3.0 (Jun 12, 08:55) → **1.5, i.e. OFF** (Jun 12, 09:57, operator call). This weekend is the first ever traded on the late window — every weekend trade is tagged (lead size, ask, duration) for a Monday judgment of whether late entry neutralizes the historical weekend bleed (–\$148 one book, –\$112 another, both pre-window).

3 · Execution safety — floor breaches & auto-scratch

Three-layer fill floor (Jun 11, ~02:40): book freshness stamps on every WSS tick; pre-sign gate ($0.50 \leq \text{best ask} \leq 0.75$ AND book age $\leq 1.5\text{s}$); last-instant re-check between signing and POST. Eliminates all floor breaches except the physically unclosable ~100–400ms wire gap.

Auto-scratch (Jun 12, 01:48): when a FOK still fills below 0.50 inside that wire gap, the bot immediately FOK-sells the entire position back (limit = fill – 0.15). Full exit → recorded as 'scratched' with round-trip cost (~\$0.30–1.50 vs a \$10 coin flip), deliberately invisible to the ladder and resolver. Sell kills or errors → holds as before (fails safe, no partial states). Evidence that breach fills are poison: they only print cheap because the book repriced against the side inside the wire gap — the 0.300 breach fill resolved LOST; 5 breaches total in ~800 trades.

4 · Resolution & ladder

Instant Polymarket-sourced verdicts (Jun 11, prov v2): ~2s after each window closes, the bot reads Polymarket's own CLOB midpoint (≥ 0.95 = won, ≤ 0.05 = lost) and steps the ladder immediately — no more 5–7 minute lag causing repeated stakes at the wrong rung. Polymarket is the only resolution source; on-chain settlement remains authoritative for money and guardrails.

Duration-weighted ladder (Jun 11, ~02:45): the \$10→\$15→\$22.50 cycle is capped at the \$15 rung on 5m markets (PAROLI_5M_MAX_STEP=1) — own-book data showed \$22.50-on-5m ran 53.2% WR / –\$224.66

over 47 trades. 15m keeps the full ladder. Wins climb, three wins cycle back to base, any loss resets — verified cycling correctly multiple times live. Scratched trades neither step nor reset the ladder.

5 · Data collection (for the June 18 review)

Exit-liquidity logger (Jun 11, ~07:08): the WSS feed's bid side — previously discarded — is now captured, and every held position's final 35 seconds are snapshotted at 1Hz (best bid, top-5 bid depth, level count, live lead). This is the dataset that decides whether an exit kill-switch (sell when the lead flips with seconds left) is viable — verdict June 18, pooled across team books. Collection paused only during the Jun 12 bot-stop window (12:53–13:05 UTC).

6 · Display truth — terminal & dashboard

Live banner (Jun 12, 08:52 + 08:59): the 30-second status line now prints the ACTIVE entry bar (weekend-aware), day PnL, balance, and all-time win rate — every value computed from the database and the same functions the entry check uses.

Dashboard de-lied (Jun 12, 09:15): the state file fed the dashboard static constants (showed 'min move 1.5bps / stake \$10' while the bot was actually holding a 3.0bps bar and a \$15 ladder stake). It now writes the live effective values. Rule reinforced in the team playbook: displays must call the same function the entry logic calls — twice now a display lied while the bot told the truth.

Current config snapshot (live right now)

Setting	Value	Setting	Value
Fire window	7–14s before close	Entry bar	1.5bps (weekend bar off)
Ask band	0.50 – 0.75	Fill floor	0.50 + auto-scratch
Book depth gate	>= \$100 in-band	Ladder	\$10/\$15 (5m cap), full on 15m
Balance floor	\$240 (real-dollar)	Daily loss limit	\$150
Loss streak halt	10 in a row	Blackout	21:00–23:00 UTC
Resolution	Polymarket only	Orders	FOK — no partial fills

Scoreboard since the sweet-spot deploy (Jun 11 13:05 → Jun 12 13:45 UTC)

19 trades: **15W–4L (79%), +\$35.54** · ladder cycled cleanly 3× · depth gate blocked 10+ phantom books (fills \$4–\$92 behind the ask) · double-side guard blocked 3 opposite-side entries · 1 wire-race breach occurred pre-scratch-deploy (won, luck) · Jun 11 closed GREEN (+\$5.75) after a –\$40 morning under the old window. Every change has a timestamped backup on the server; rollback for any item is a one-line file restore + restart.

Ghost In the System — internal team document. No keys, wallets, or endpoints included.